


# Big Data – Introduction to Apache™ Hadoop®



Sridhar Paladugu

<http://www.linkedin.com/in/sridharpaladugu/>

# Agenda

---



- **Big Data Concept**
- **What is Hadoop**
- **Architecture of Hadoop**
- **Basic Hadoop maintenance tasks**
- **Hadoop Sample Application**
- **Hadoop Eco System**
- **Pig Overview and sample application**
- **Hive Overview and sample application**
- **Hbase Overview**
- **Conclusion**
- **Questions**

# **BIG DATA is often defined by 4 V's**

## **Volume**

Terabytes or even petabytes of data for processing and ever increasing yet



## **Velocity**

Pace at which the data are to be consumed.



## **Variety**

Structured and unstructured data

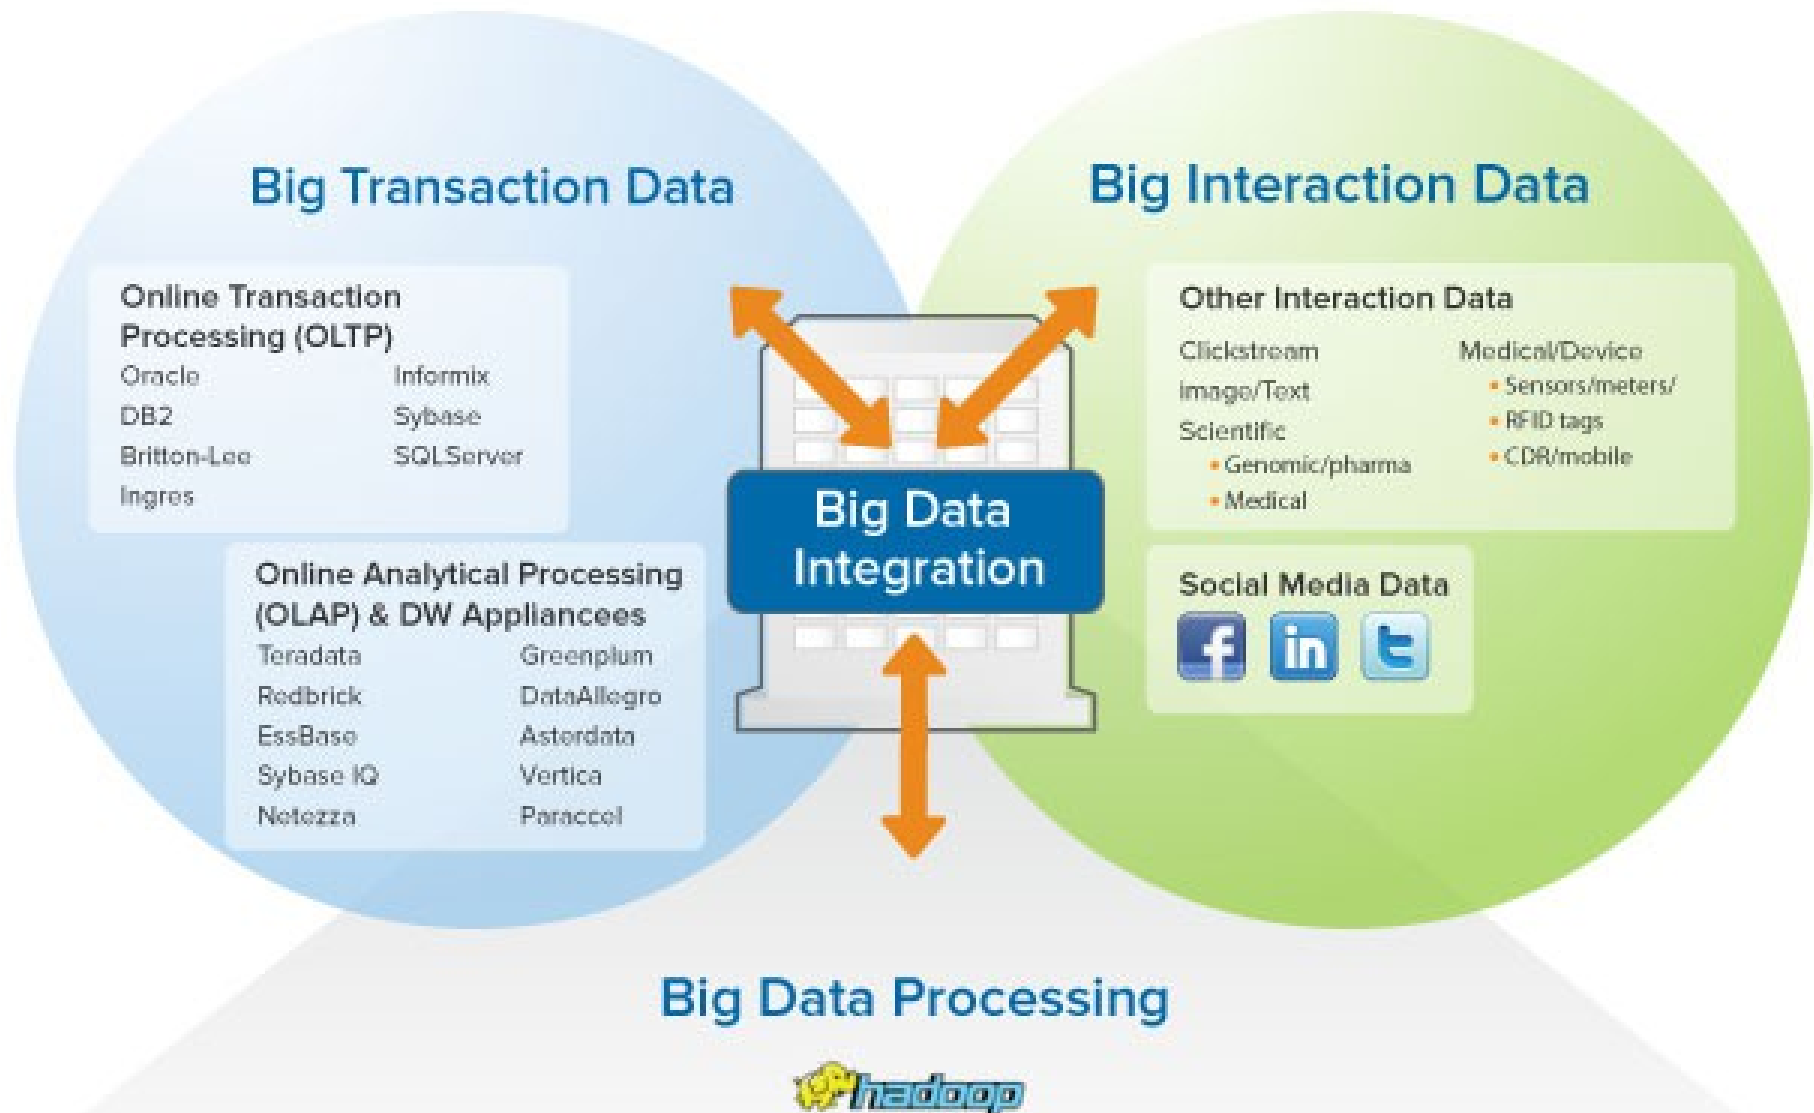


## **Veracity**

Data inconsistency and incompleteness




**Definition:** Big data is the confluence of the three trends consisting of Big Transaction Data, Big Interaction Data and Big Data Processing



# What is Hadoop?

---




- open source project for big data processing.
- provides an infrastructure for a distributed storage and a distributed computation.
- Highly Scalable
- High degree of fault tolerant.
- Cost effective
- Flexible data format support.
- Redundant and reliable.



# Core Components of Hadoop

- Distributed File System (HDFS)
- Parallel computing framework (Map/Reduce)
- Hadoop commons - Job scheduler and Management features



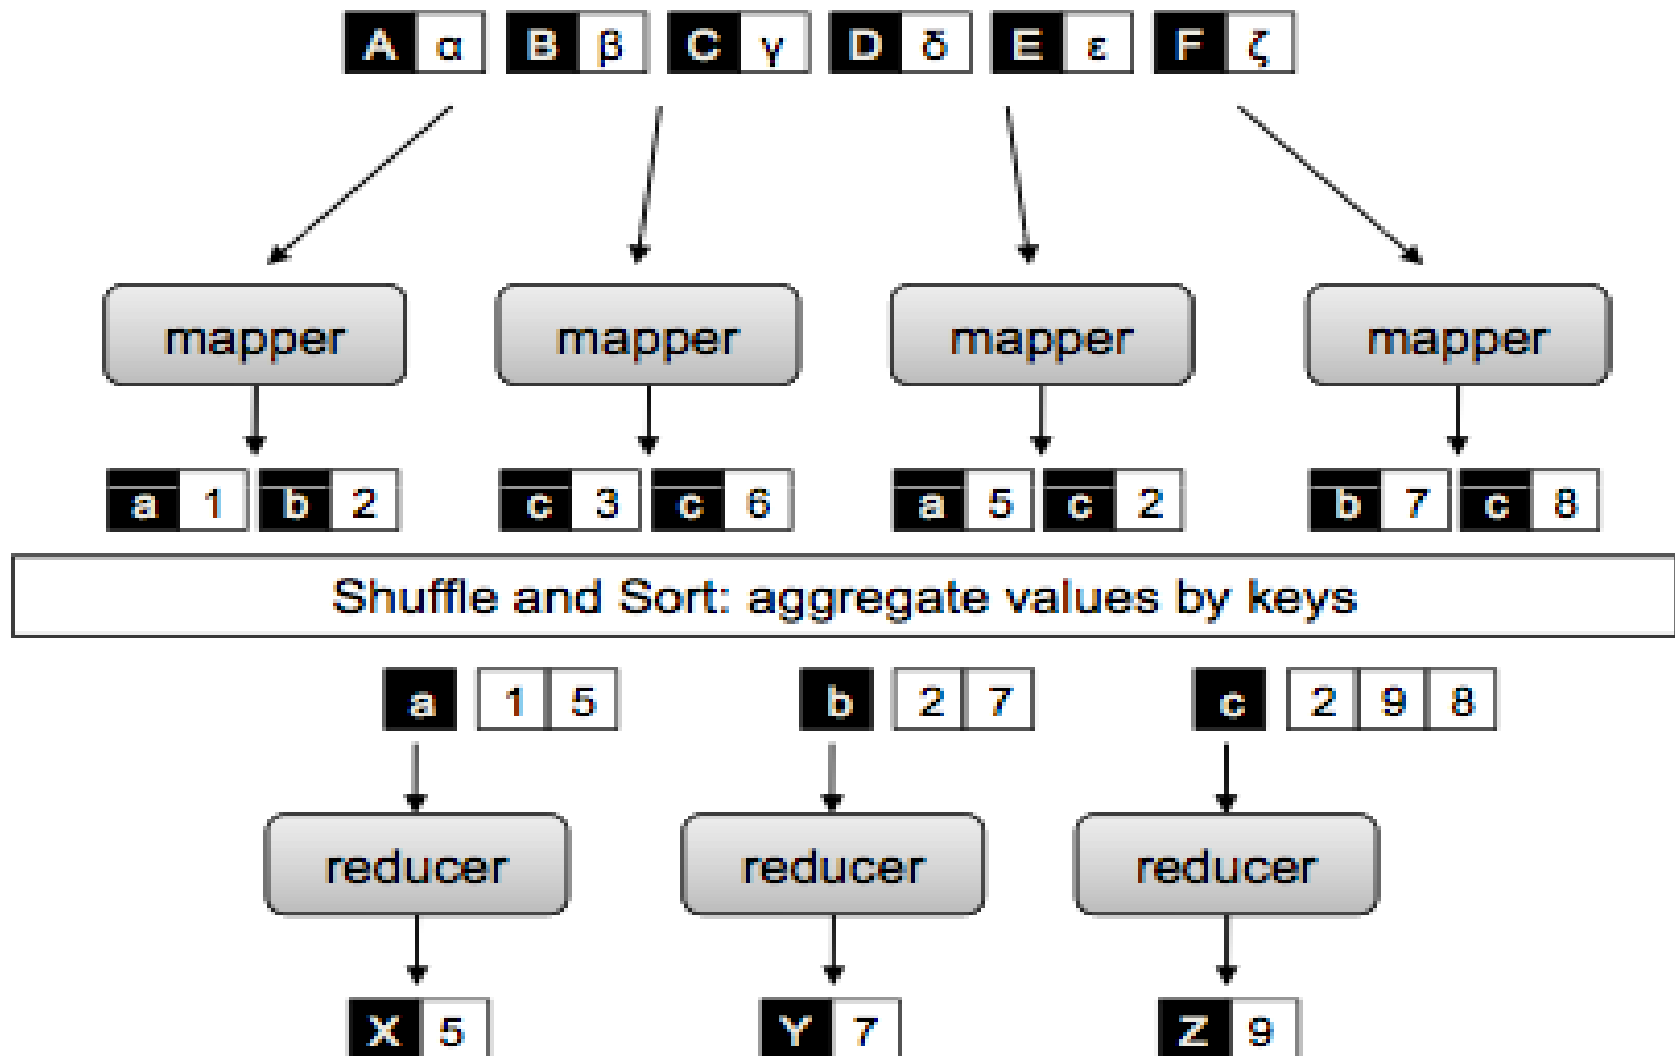


# MapReduce

- Programming model designed to process large data sets.
- Input & Output: each a set of key/value pairs
- **Map**
  - ▢ Processes input key/value pair and Produces set of key/value pairs
- **Reduce**
  - ▢ Shuffle map output and Combines all intermediate values for a each key.
  - ▢ Merge and produce output.



# Map Reduce







# HDFS

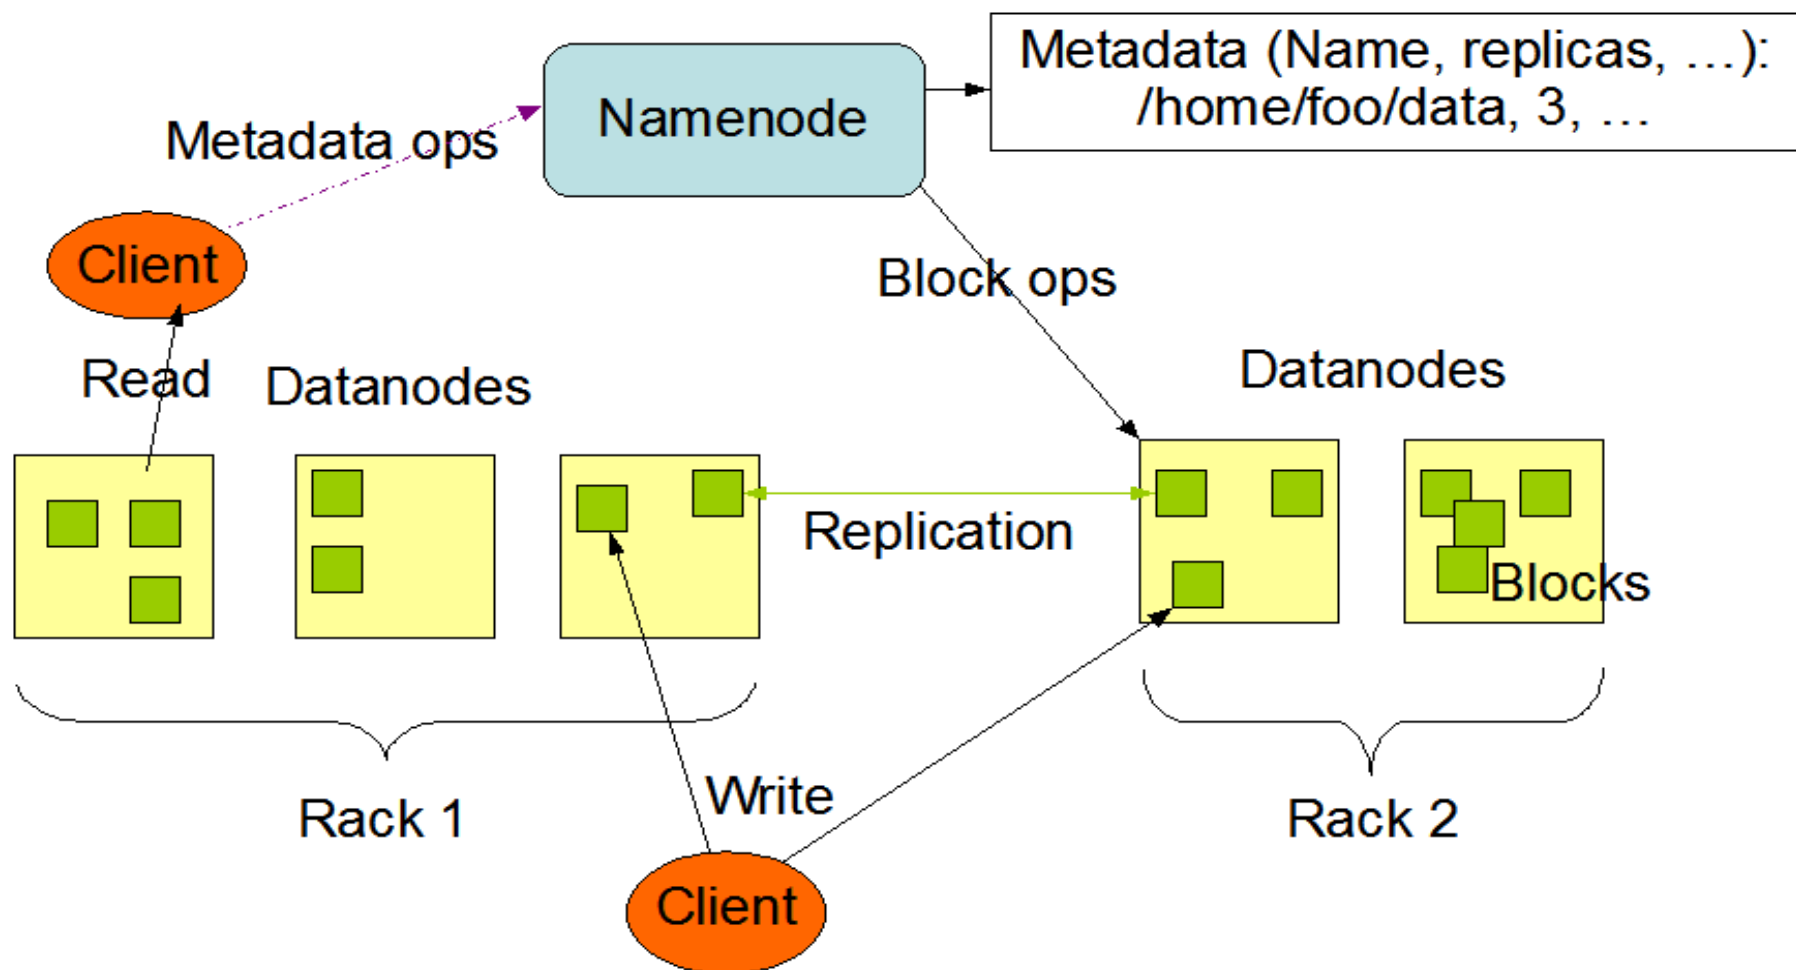
---

- ♦ **Distributed File system**
- ♦ **Write-once, Read-many model**
- ♦ **Data access via MapReduce streaming**
- ♦ **Fault Tolerance - Automatically and seamlessly recover from failures**
- ♦ **Load Balancing - Place data intelligently for maximum efficiency and utilization.**
- ♦ **Tunable Replication - Multiple copies of each file provide data protection and computational**
- ♦ **Security - POSIX-based file permissions for users and groups with optional LDAP integration**

# HDFS



## HDFS Architecture





# Working with HDFS

- Using Linux Shell we can run hdfs file system commands.
  - ▢ **\$hadoop fs -ls /user/sridhar**
- Storing a file to HDFS
  - ▢ **\$hadoop fs -copyFromLocal myfile.txt /user/sridhar/volume1**
- Retrieving a file from HDFS
  - ▢ **\$hadoop fs -copyToLocal /user/sridhar/volume1 myfile.txt**
- Copy file to HDFS replication
  - ▢ **\$hadoop distcp hdfs://namenodeA/user/sridhar hdfs://namenodeB/user/sridhar**
- Merge output
  - ▢ **\$hadoop fs -getmerge /user/sridhar/equityPricing eqtyPricing\_11302012.bcp**



# Working with HDFS

- Using Linux Shell we can run hdfs file system commands.
  - ▢ **\$hadoop fs -ls /usr/sridhar**
- Storing a file to HDFS
  - ▢ **\$hadoop fs -copyFromLocal myfile.txt /usr/sridhar/volume1**
- Retrieving a file from HDFS
  - ▢ **\$hadoop fs -copyToLocal /usr/sridhar/volume1 myfile.txt**
- Copy file to HDFS replication
  - ▢ **\$hadoop distcp hdfs://namenodeA/usr/sridhar hdfs://namenodeB/usr/sridhar**
- Merge output
  - ▢ **\$hadoop fs -getmerge /usr/sridhar/equityPricing eqtyPricing\_11302012.bcp**

# Working with HDFS : Sqoop

---

- **Open source tool to transfer data in and out from hadoop tp RDBMS.**
- **Utilizes MapReduce to do parallel execution and fault tolerance.**
- **Support Hive and hbase.**
- **Import from table or query**
  - **Divide table in to range using primary key or specified column.**
  - **Create mappers for each range and executes in parallel.**
  - **Can output in all formats supported by hdfs.**
  - **Generates java class for generated hdfs file.**
- **Export to Database**
  - **Utilize MapReduce to read data concurrently.**
  - **Utilize bulk insert using Parallelism.**

# Import and Export using Sqoop

- Import using query

- ▢ `$sqoop import`

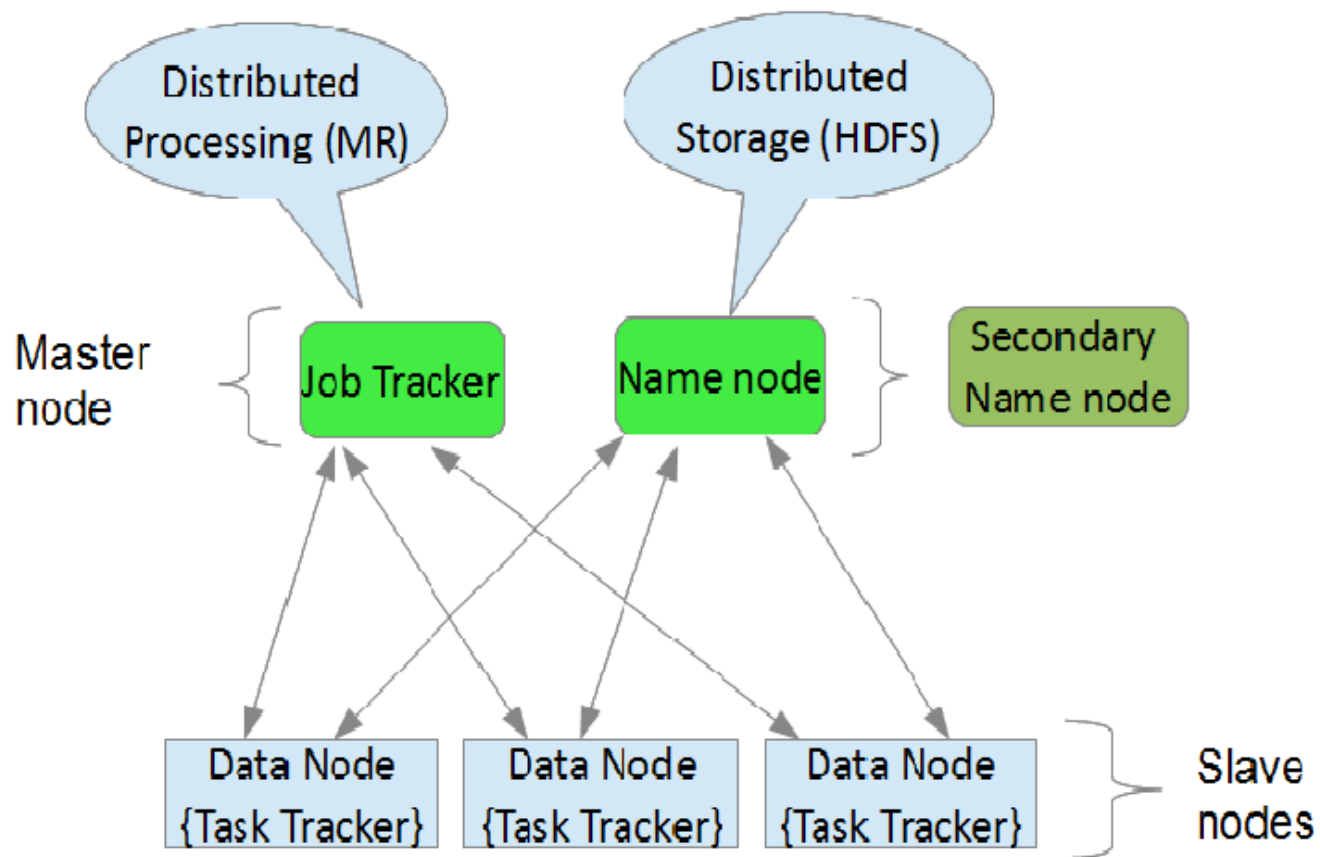
- `--connect jdbc:oracle:thin:@tradeNode-1:1521:eqty`
  - `--username eqtyTrader --password passw0rd`
  - `--as -textfile`
  - `--query 'select t.tradeId,b.bookName, cp.cunterparty, t.tradeDate,t.mtm'\`  
`' from Trade t '\`  
`' join Counterparty cp on t.counterpartyFk=cp.id '\`  
`' join eodMarks20121130 m on t.tradeFk = m.tradeFk '\`
  - `--split-by t.tradeId`
  - `--fields-terminated-by '\t' --lines-terminated-by '\n' \`
  - `--targetDir /usr/eqty/mtm/20121130`
  - ▢ *[Incremental, last-value, hive-import,hive-overwrite, options-file ]*

- Export to database

- `$sqoop export`

- `--connect jdbc:oracle:thin:@tradeNode-1:1521:eqty`
  - `--username eqtyTrader --password passw0rd`
  - `--table EOD_MTMSummary`
  - `--update-key id --update-mode=allowinsert`
  - `--export-dir /path/to/data`

# Hadoop Cluster view





# Hadoop : Basic Configuration

Hadoop is mainly configured using files `core-site.xml` , `hdfs-site.xml` , `mapred-site.xml`

**Core-site.xml;**            **specifies the name node related configuration.**

```
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<configuration>
  <property>
    <name>hadoop.tmp.dir</name>
    <value>/usr/local/hadoop-1.0.4/tmp</value>
  </property>
  <property>
    <name>fs.default.name</name>
    <value>hdfs://hdp-master.localdomain.com:54310</value>
  </property>
</configuration>
```






# Hadoop : Basic Configuration

- **Hdfs-site.xml** specify information related to hdfs

```
<?xml version="1.0"?>
<configuration>
  <property>
    <name>dfs.name.dir</name>
    <value>/usr/local/hadoop-1.0.4/hdfs/name</value>
  </property>
  <property>
    <name>dfs.data.dir</name>
    <value>/usr/local/hadoop-1.0.4/hdfs/data</value>
  </property>
  <property>
    <name>dfs.replication</name>
    <value>3</value>
  </property>
</configuration>
```




# Hadoop : Basic Configuration

- **Mapred-site.xml** configuration related to job tracker is specified here

```
<?xml version="1.0"?>
<configuration>
<property>
  <name>mapred.job.tracker</name>
  <value>hdp-master.localdomain.com:54311</value>
  <description>job tracker address</description>
</property>
<property>
  <name>dfs.hosts.exclude</name>
  <value>/path/to/hadoop/dfs_excludes</value>
  <final>true</final>
</property>
<property>
  <name>mapred.hosts.exclude</name>
  <value>/path/to/hadoop/mapred_excludes </value>
  <final>true</final>
</property>
<property><name>mapred.reduce.parallel.copies</name> <value>5<value> </property>
```

# Hadoop : Basic maintenance tasks



- Start Hadoop
  - ▣ Standalone mode: run as a single process on a single node.
  - ▣ Pseudo-distributed mode: run all services in separate processes on a single node.
  - ▣ Fully-distributed mode: run all services in separate processes across multiple nodes.
    - **sridhar@hdp-standalone:/usr/local/hadoop/bin/start-all.sh**  
**sridhar@hdp-standalone:/usr/local/hadoop-1.0.4/myscripts\$ jps**  
**15020 Jps**  
**24033 NameNode**  
**24815 TaskTracker**  
**24274 DataNode**  
**24504 SecondaryNameNode**  
**24589 JobTracker**  
**9359 RunJar**
- Stop Hadoop
  - **\$/usr/local/hadoop/bin/stop-all.sh**



# Hadoop : Basic maintenance tasks

- Add nodes to cluster
  - ▢ Setup new linux machine and install the distribution; hdp-node-3  
sridhar@hdp-master:/usr/local/hadoop-1.0.4/conf\$ vi slaves  
hdp-node-1.localdomain.com  
hdp-node-2.localdomain.com  
**hdp-node-3.localdomain.com**
  - ▢ In new slave machine, run below commands  
sridhar@hdp-node-3:/usr/local/hadoop-1.0.4/bin/start datanode  
sridhar@hdp-node-3:/usr/local/hadoop-1.0.4/bin/start tasktracker
  - ▢ In Master node, run  
sridhar@hdp-master:/usr/local/hadoop-1.0.4/bin/start-balancer.sh



# Hadoop : Basic maintenance tasks

---

- **Decommission a node**

```
sridhar@hdp-master:/usr/local/hadoop-1.0.4/conf$vi dfs_excludes
```

```
hdp-node-3.localdomain.com
```

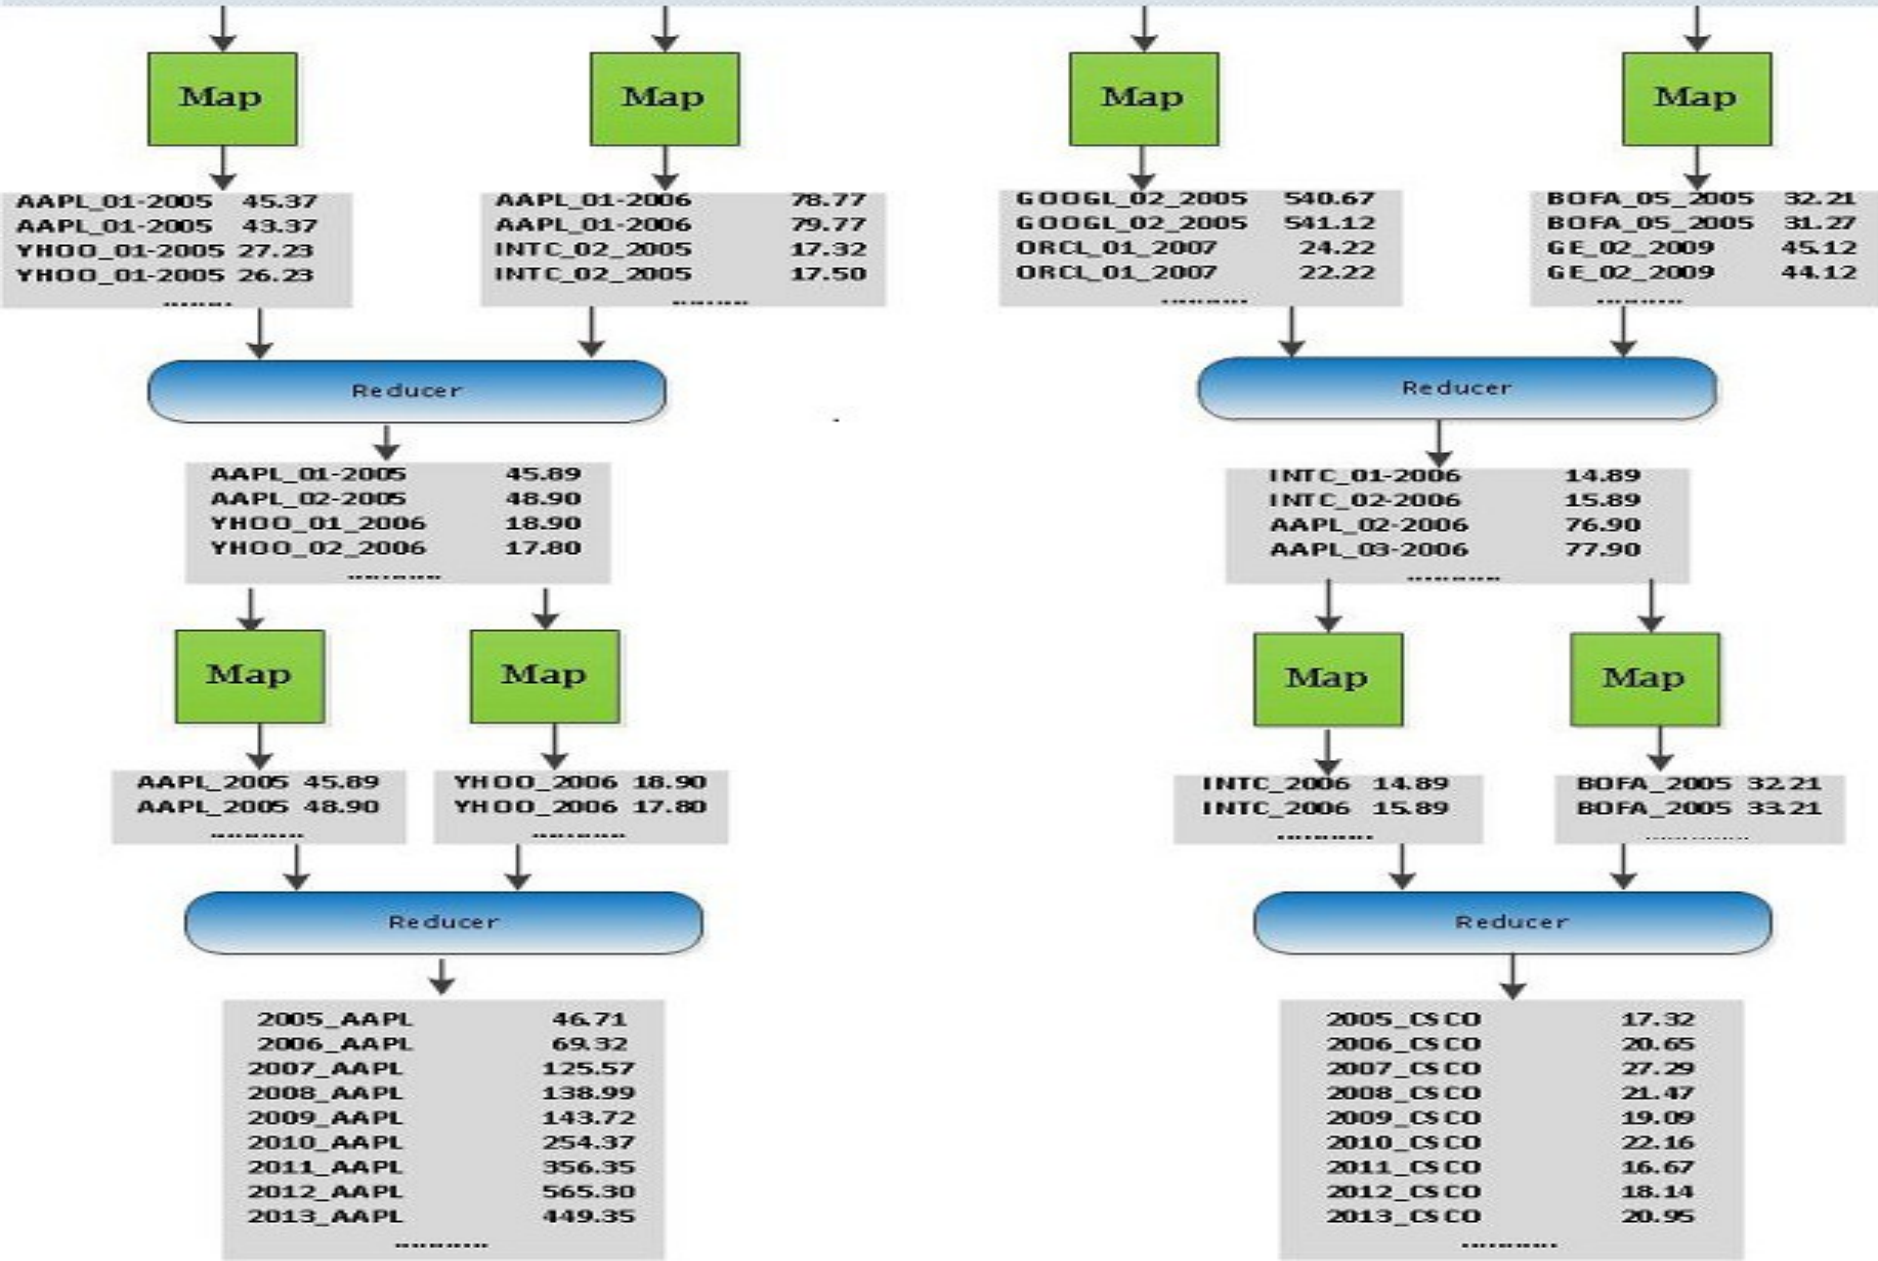
```
sridhar@hdp-master:/usr/local/hadoop-1.0.4/conf$vi mapred_excludes
```


```
hdp-node-3.localdomain.com
```

```
sridhar@hdp-master:/usr/local/hadoop-1.0.4/bin$hadoop dfsadmin -refreshNodes
```

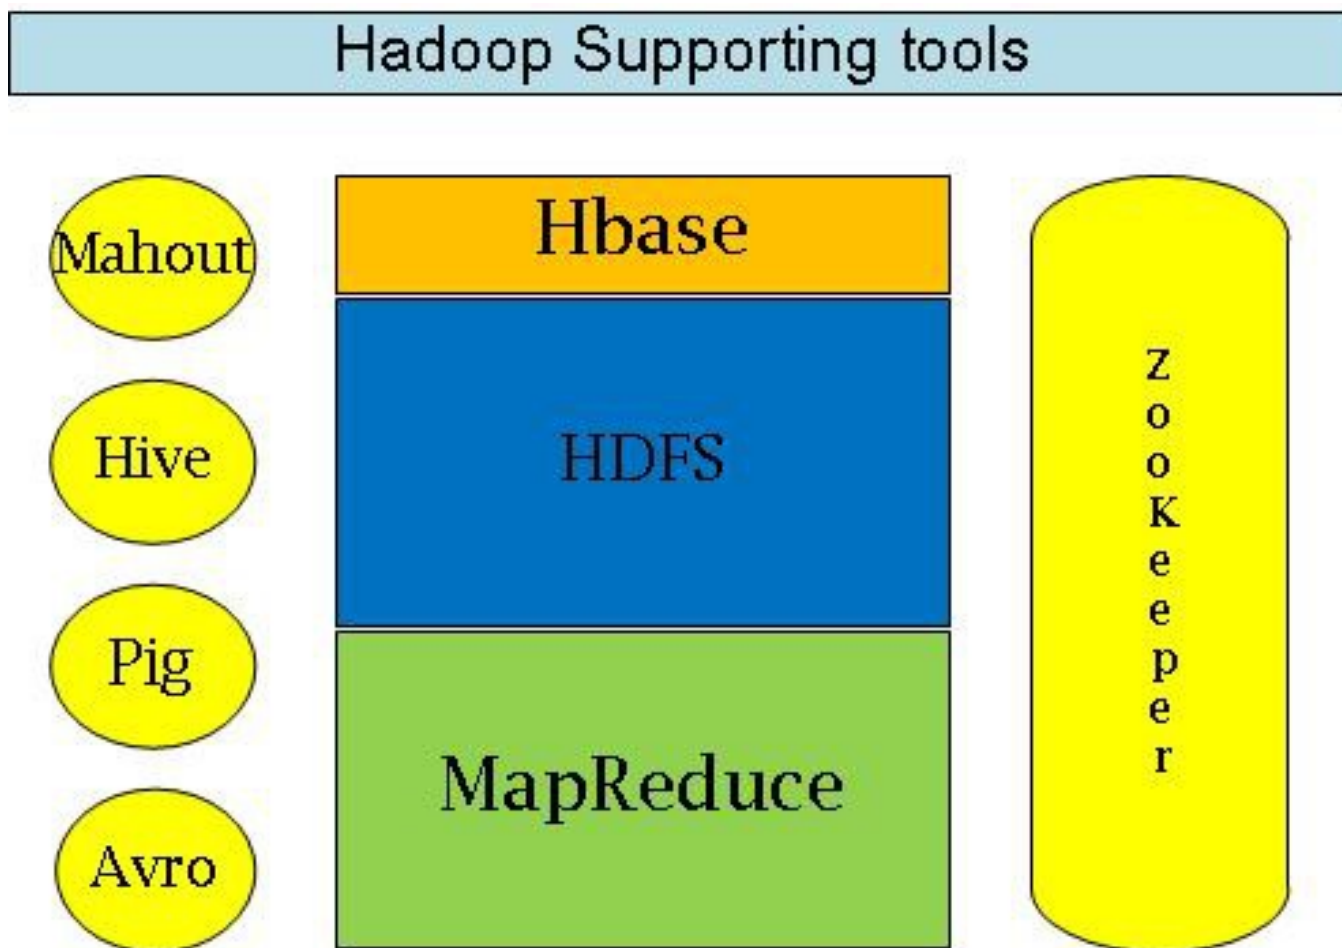
```
sridhar@hdp-master:/usr/local/hadoop-1.0.4/bin$hadoop mradmin -refreshNodes
```

TICKER	PRICE_DATE	PREV_OPEN	PREV_CLOSE	OPEN	HIGH	LOW	VOLUME	CLOSE
AAPL	01/01/2005	45	45.27	45.27	45.10	45.37	200000	45.37
YHOO	01/01/2005	27	27.27	27.27	27.27	27.01	500000	27.27
GOOGL	01/05/2013	670.90	680.16	680.16	687.31	682.10	100000	687.21






# Hadoop Eco System



# Apache PIG

- High level structured query language for analyzing data stored in Hadoop cluster.
- Provide a run-time time environment to generate Map Reduce programs from a high level data-flow statements.
- Pig Latin, the data-flow language
- Grunt pig shell, run-time
- Not a full featured programming language, or not equivalent to SQL.
- Extensible
- Why Pig?





# Apache PIG Run-time environment



- Pig run-time is started using ;

▢ `$/usr/local/pig/pig`

**\$grunt>**

**\$grunt>** data = LOAD(...) AS(...)

**\$grunt>** filter\_by\_v = FILTER data by (vendor = 'xyz')

**\$grunt>** group\_by\_dt = GROUP filter\_by\_v BY (shop\_date)

**\$grunt>** transactions = FOR EACH group\_by\_dt GENERATE

group, sum(amt)


**\$grunt>** STORE group\_by\_v INTO ( .... )

# Apache Pig – Development and Testing




- Pig run time supports User defined functions defined in Java, Ruby, Python code
- Running a pig script
  - ▢ **`$/usr/local/pig/pig -x sampleScript_1.pig`**
- Describe -- print the schema of a data structure
  - **DESCRIBE data\_alias;**
- Explain – print how map reduce jobs are being generated for a script
  - **Explain data\_alis**
  - ▢ **`/usr/local/pig/pig -x local -e 'explain -script sampleScript_1.pig'`**
- Illustrate --takes a sample of your data and runs it through your script
  - ▢ **`/usr/local/pig/pig -x local -e 'illustrate -script illustrate.pig'.`**

Transaction Date	Description	Amount	Comments
07/21/2012	TARGET 00010983 NAPERVILLE, IL	5.11	Return processed***
07/20/2012	PAYMENT	-1330.33	
07/20/2012	TRU HOLIDAY EXPRESSQPS NAPERVILLE, IL	10.74	
07/13/2012	TARGET 00010983 NAPERVILLE, IL	6.27	
07/10/2012	TARGET 00010983 NAPERVILLE, IL	36.04	



TARGET #00010983 NAPERVILLE, IL	6.27
TARGET #0010983 NAPERVILLE, IL	36.04
.....	
.....	



PETSMART INC 54, NAPERVILLE, IL	22.06
PETSMART INC 54, NAPERVILLE, IL	11.06
.....	
.....	



TARGET #00010983 NAPERVILLE, IL	2300.67
SAMSClub #6620 NAPERVILLE, IL	1892.00
PETSMART INC 54, NAPERVILLE, IL	428.06
CVS PHARMACY #9923 AURORA, IL	321.22
.....	



# Apache PIG : Sample script

```
all_credit_tx = LOAD 'hdfs://hdp-  
master.localdomain.com:54310/user/sridhar/data/cred  
it_card_tx.txt' AS (transactionDate:chararray,  
transactionDesc:chararray, amount:float);  
  
filter_returns = FILTER all_credit_tx BY  
(transactionDesc != 'PAYMENT');  
  
group_by_vendor = GROUP filter_returns by  
transactionDesc;  
  
sum_vendor_exp = FOREACH group_by_vendor GENERATE  
group, SUM(filter_returns.amount);  
  
STORE sum_vendor_exp INTO  
'/usr/sridhar/pig/out/SPENDING_ANALYSIS';
```

# Apache PIG : Sample script Results



Run the script using

```
$pig -f ExpByVendor.pig
```

```
SKECHERS-USA #119    CHICAGO, IL    2127021.75
```

```
SIX FLAGS GREAT AMERICA IL    IL    2127777.10
```

```
76 10115095          NAPERVILLE, IL    2122521.27
```

```
BABIES R US #6447  QPS NAPERVILLE, IL  2126748.77
```



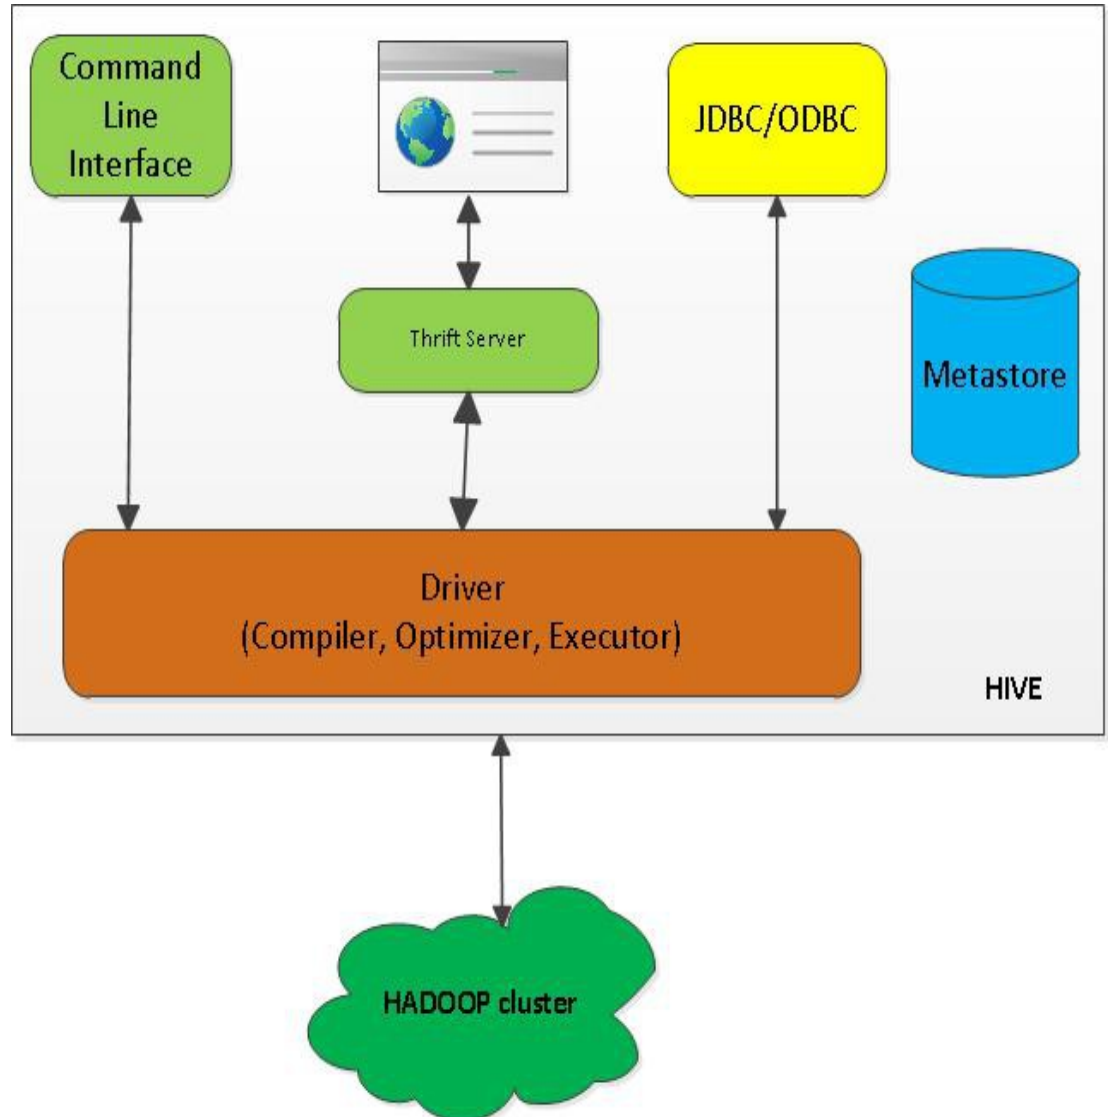
# Apache PIG: Conclusion

- provides easy way of implementing Map reduce jobs.
- Provides extensible features via UDF
- Provides less complex way of creating job flows
- Provides Lazy evaluation
- Ability to store data anytime during flow.
- Provides out of box performance enhancers



# Apache HIVE

- A distributed data warehouse engine built on top of Hadoop
- Provides HQL interface to HDFS data
- Utilize Map-Reduce to execute queries
- Work with data in various formats
- Extensible
- Supports Partitions and Joins





# Apache HIVE : Data Types

---

- Domain objects
  - ▢ Tables, partitions, Buckets
- Data Types
  - ▢ Primitives
    - TINYINT, SMALLINT, INT, BIGINT, BOOLEAN, FLOAT, DOUBLE, STRING
  - ▢ Complex types
    - STRUCT, MAP, ARRAY



# Apache HIVE – DDL operations



- Start hive CLI
  - ▢ \$hive
- **hive>DROP TABLE IF EXISTS creditcard\_transactions;**
- **hive>CREATE EXTERNAL TABLE creditcard\_transactions (**  
**transactionDate STRING, description STRING, amount FLOAT, comments**  
**STRING )**  
**ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LINES TERMINATED BY**  
**'\n'**  
**LOCATION '/usr/sridhar/creditcard\_txs/';**
- **hive>show tables;**
- **hive> describe creditcard\_transactions;**



# Apache HIVE : DML statements

- **Selecting data**

```
select * from creditcard_transactions
where UNIX_TIMESTAMP(transactionDate,'yyyy-MM-dd') >=
      UNIX_TIMESTAMP('2012-01-01', 'yyyy-MM-dd')
and UNIX_TIMESTAMP(transactionDate,'yyyy-MM-dd') <=
      UNIX_TIMESTAMP('2012-01-30', 'yyyy-MM-dd')
```

- **Supports Group BY, Having and Limit**

```
  Select description, sum(amount)
  from creditcard_transactions
  where transactionDate='2012-01-02'
  group by description
  having sum(amount) < 3000
  sort by description desc LIMIT 10;
```

- **Writing data out to table or file system**

```
  INSERT INTO OVERWRITE saleBy_Vendaor SELECT ...
  INSERT INTO OVERWRITE [LOCAL] directory path_to_write SELECT...
```



# Apache HIVE: Table Partitioning

- Hive organizes table as partitions.
- each partition as a separate file
- Data defined in one partition belong to a key, making access fast
- Partitions can have sub-partitions.
- While loading data into partition table we need to specify partition.
- Dynamic partitioning is allowed\*

```
CREATE TABLE mgd_creditcard_transactions (  
    description STRING, amount FLOAT, comments STRING)  
PARTITIONED BY(transactionDate STRING);
```



# Apache HIVE: Populating a Partition

- Exclusive partition load

```
INSERT OVERWRITE TABLE mgd_creditcard_transactions  
PARTITION (transactionDate='2012-01-02')  
SELECT description,amount,comments,transactionDate  
FROM creditcard_transactions  
WHERE transactionDate='2012-01-02';
```

- Dynamic partition load

```
SET hive.exec.dynamic.partition = true;  
SET hive.exec.dynamic.partition.mode = nonstrict;  
SET hive.exec.max.dynamic.partitions.pernode = 200;  
set hive.exec.max.dynamic.partitions = 200000;
```

```
INSERT OVERWRITE TABLE mgd_creditcard_transactions  
PARTITION (transactionDate)  
SELECT description,amount,comments,transactionDate  
FROM creditcard_transactions;
```



# Apache HIVE: Extensibility

- Support custom code via UDF.

```
Package com.sridhar.hadoop.samples.hive;  
import java.math.BigDecimal;  
import java.text.NumberFormat;  
import org.apache.hadoop.hive.ql.exec.UDF;  
import org.apache.hadoop.io.FloatWritable;  
import org.apache.hadoop.io.Text;  
public class CurrencyFormatter extends UDF {  
    public Text evaluate(final Text input) {  
        if (input == null) { return null; }  
        BigDecimal d = new BigDecimal(input.toString());  
        return new Text(NumberFormat.getCurrencyInstance().format(d));  
    }  
}
```

- **Register the function**

```
CREATE FUNCTION CU_FMT AS com.sridhar.hadoop.samples.hive.CurrencyFormatter
```

- **Use in hql**

```
select transactionDate, Description, CU_FMT(amount)  
from mgd_creditcard_transactions where transactionDate='2012-01-02';
```

# Apache HIVE : Conclusion

---

- **Advantages**

- ▢ Provides easy interface to HDFS by hiding complexity of Map Reduce.
- ▢ Support SQL-based queries and thus less learning curve.
- ▢ provides External Table feature to leverage to keep data in HDFS
- ▢ Provides meta store which allows decoupling application logic from physical data structure
- ▢ Support for data partitioning for faster access
- ▢ Easy to plug-in custom implementation for user defined functions

- **Disadvantages**

Not designed for OLTP

Latency issues because of need to generate Map Reduce jobs



# Apache HBase

---

Implemented using Google's Big Table paper.

A column-oriented database stores data in key value pairs.

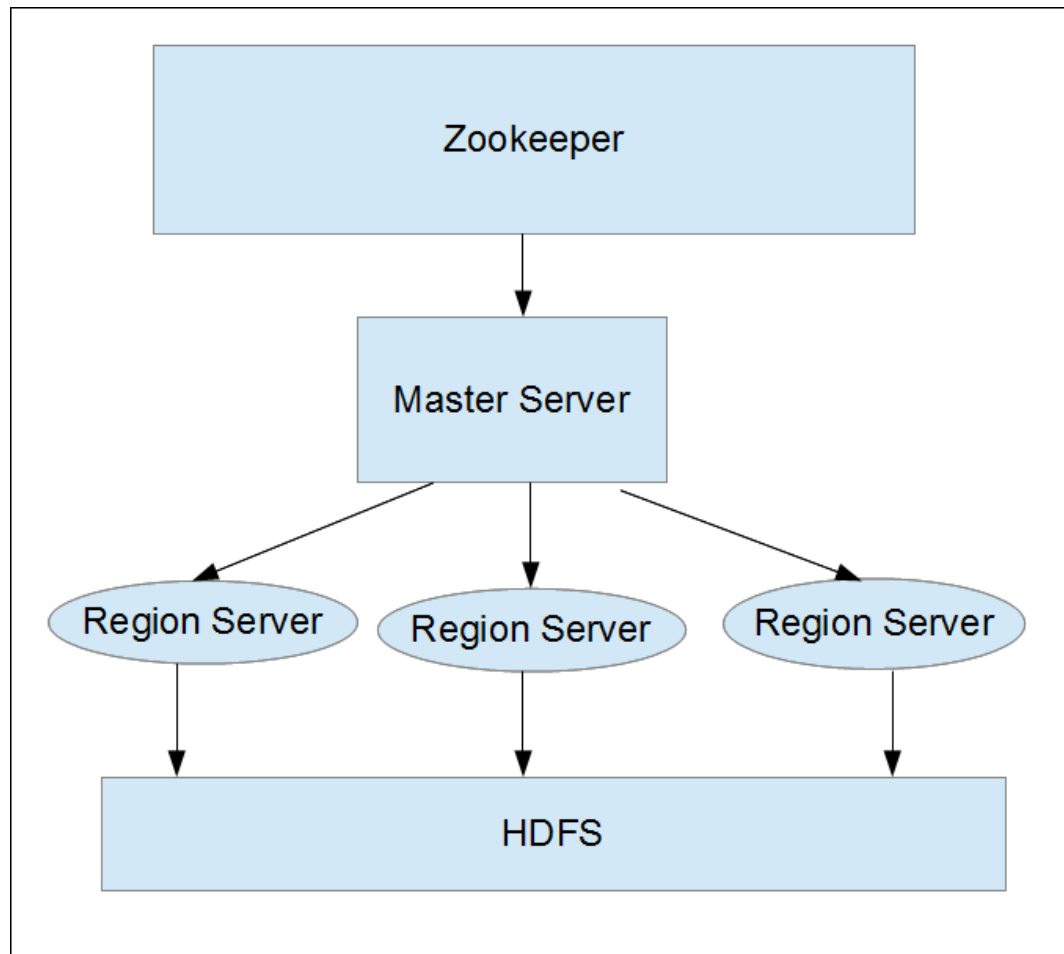
## **Characteristics:**

- Sparse
- Distributed
- Sorted
- Multidimensional
- Version aware

## **Features:**

- Vertical and horizontally scalable
- Automatic Fail-over
- Consistent reads and Writes
- Integration with Map Reduce framework
- Simple Java API for Client access

# Apache HBase - Architecture







# Apache Hbase – Data Model

- Data is stored in Table
- Row = {key, value}
  - ▢ Key can be primitive type or custom data structure
  - ▢ Value is typically a set of column families
- A Column family is a map of column cells
  - ▢ row\_1 {col\_family {cell1, cell2....} }
- Column cell is a key value pair with version
  - ▢ Row\_1 {col\_family\_1: {cell\_1, LIST{cell\_value, cell\_version}  
col\_family\_1: {cell\_2, LIST{cell\_value, cell\_version}  
..... }  
.....}
- Java construct:  
  
SortedMap<Key, List<SortedMap<key, List<Value, timestamp>>>



# Apache Hbase – Table Storage

## Logical View:

- Table is made of regions.
- Region is a range of Keys, key range is inclusive.
- Each Region Server hosts one or more regions
- Regions grow dynamically as data grows

## Physical View:

- WAL (Write Ahead Log)
- MemStore – in memory storage
- HFile – an indexed file. Each flush operation creates one file
- -ROOT-
- .META – Region meta-base



# Apache HBase : Hbase Shell

Login to Shell

```
sridhar@hdp-master:/usr/local/hbase-0.94.8/bin$ hbase shell
```

Create Table

```
hbase(main):006:0> create 'books', {NAME => "info", VERSIONS => 5}
```

Insert data to hbase table using put

```
hbase(main):006:0> put 'books', 'book2', 'info:title','HBase In Action'
```

Fetch data:

```
hbase(main):007:0> scan 'books'
```

ROW	COLUMN+CELL
-----	-------------

book1	column=info:title, timestamp=1371505821795, value=Hadoop Definitive guide 2nd version
-------	--

book2	column=info:title, timestamp=1371525002008, value=HBase InAction
-------	--

2 row(s) in 0.1590 seconds

Fetch Data for a specific cell value:

```
hbase(main):002:0> scan 'books', { COLUMNS => 'info:title', LIMIT => 10, FILTER => "ValueFilter( =,  
'binaryprefix:Hadoop' )" }
```

ROW	COLUMN+CELL
-----	-------------

book1	column=info:title, timestamp=1371505821795, value=Hadoop Definitive guide 2nd version
-------	---

1 row(s) in 0.1550 seconds

Find rows in a table:

```
hbase(main):003:0> count 'books'
```

2 row(s) in 0.0140 seconds



# Apache HBase: Import/Export Data

- Export/Import MapReduce jobs

Export [-D<property=value>]\* <tablename> <outputdir> [<versions>  
[<starttime> [<endtime>]] [^[regex pattern] or [Prefix] to filter]]

```
❏ sridhar@hdp-master:hbase org.apache.hadoop.hbase.mapreduce.Export  
-Dmapred.output.compress=true  
-Dmapred.output.compression.codec=org.apache.hadoop.io.compress.Gz  
ipCodec  
-Dmapred.output.compression.type=BLOCK 'books'  
/home/sridhar/Downloads/hbaseExport
```

Import <tablename> <inputdir>

```
❏ sridhar@hdp-master:hbase org.apache.hadoop.hbase.mapreduce.Import  
'books_restored' /home/sridhar/Downloads/hbaseExport
```

# Apache HBase: Import Data



- Using ImportTsv:

```
create 'daily_stock_prices', {NAME => "price_info", VERSIONS => 5}
```

```
sridhar@hdp-master:/usr/local/hbase-0.94.8/bin$ hbase
```

```
org.apache.hadoop.hbase.mapreduce.ImportTsv
```

```
-Dimporttsv.columns=HBASE_ROW_KEY,price_info:exchange,price_info:ticker,price_info:priceDate,price_info:open,price_info:high,price_info:low,price_info:close,price_info:volume,price_info::adjustedClose daily_stock_prices /user/sridhar/data/nasdaq.csv
```

```
hbase(main):002:0> get 'daily_stock_prices', 'NASDAQ_ACAT_2006-10-05'
```

```
COLUMN
```

```
CELL
```

```
price_info::adjustedClose timestamp=1371535356480, value=16.18
```

```
price_info:close timestamp=1371535356480, value=17.09
```

```
price_info:exchange timestamp=1371535356480, value=NASDAQ
```

```
price_info:high timestamp=1371535356480, value=17.47
```

```
price_info:low timestamp=1371535356480, value=16.94
```

```
price_info:open timestamp=1371535356480, value=17.4
```

```
price_info:priceDate timestamp=1371535356480, value=2006-10-05
```

```
price_info:ticker timestamp=1371535356480, value=ACAT
```

```
price_info:volume timestamp=1371535356480, value=120800
```

```
9 row(s) in 0.5280 seconds
```



# Apache HBase: Conclusion

- HBase is not a RDBMS replacement.
- Use only when data is huge.
- Use if application does not use and RDBMS features.
- Use when data is loaded and accessed using key.
- Use when application has a variable schema.
- Use when data volumes are expected to grow and need to scale cost effectively.

# Conclusion

---

- At present Hadoop is only cost effective parallel computing environment for Big Data
- Hadoop is not a replacement for current database systems.
- Hadoop is 10 years old and growing market share.
- Hadoop MR2 is maturing and coming up with more robust features.

# Thank you.

**Sridhar Paladugu**

**<http://www.linkedin.com/in/sridharpaladugu/>**